

Midterm Project Presentation

Abhi Desai
Vinish Kandadi

Common instructions for all models

- Input: 70,000 handwritten digit images , flattened to 784 features
- Data Preprocessing:
 - Pixel values scaled from [0-255] to [0-1]
 - Data split: 80% training, 20% testing
- Performance Metrics used :
 - accuracy
 - Precision
 - Recall
 - F1-score
 - Confusion matrix

Logistic Regression

- Overall:
Data Preparation → Data Split (80% train/20% test) → Multinomial Logistic regression (softmax regression), 1000 iterations) → Training & Prediction (Fit & Predict) → Evaluation (92% accuracy, confusion matrix, per-digit performance) → Visualization (Sample predictions & accuracy plots)
- Logistic Regression :
Model → Probability Output → Hard Classification (select highest probability class, since image must belong to exactly one digit class)

Logistic Regression - Performance


Model Performance Metrics:

Accuracy: 0.9164
Precision: 0.9162
Recall: 0.9164
F1-score: 0.9162


Detailed Performance Report:

	precision	recall	f1-score	support
0	0.95	0.97	0.96	1381
1	0.95	0.96	0.96	1575
2	0.92	0.89	0.90	1398
3	0.89	0.89	0.89	1428
4	0.93	0.91	0.92	1365
5	0.87	0.87	0.87	1263
6	0.94	0.95	0.94	1375
7	0.92	0.94	0.93	1459
8	0.90	0.88	0.89	1365
9	0.89	0.89	0.89	1391
accuracy			0.92	14000
macro avg	0.92	0.92	0.92	14000
weighted avg	0.92	0.92	0.92	14000

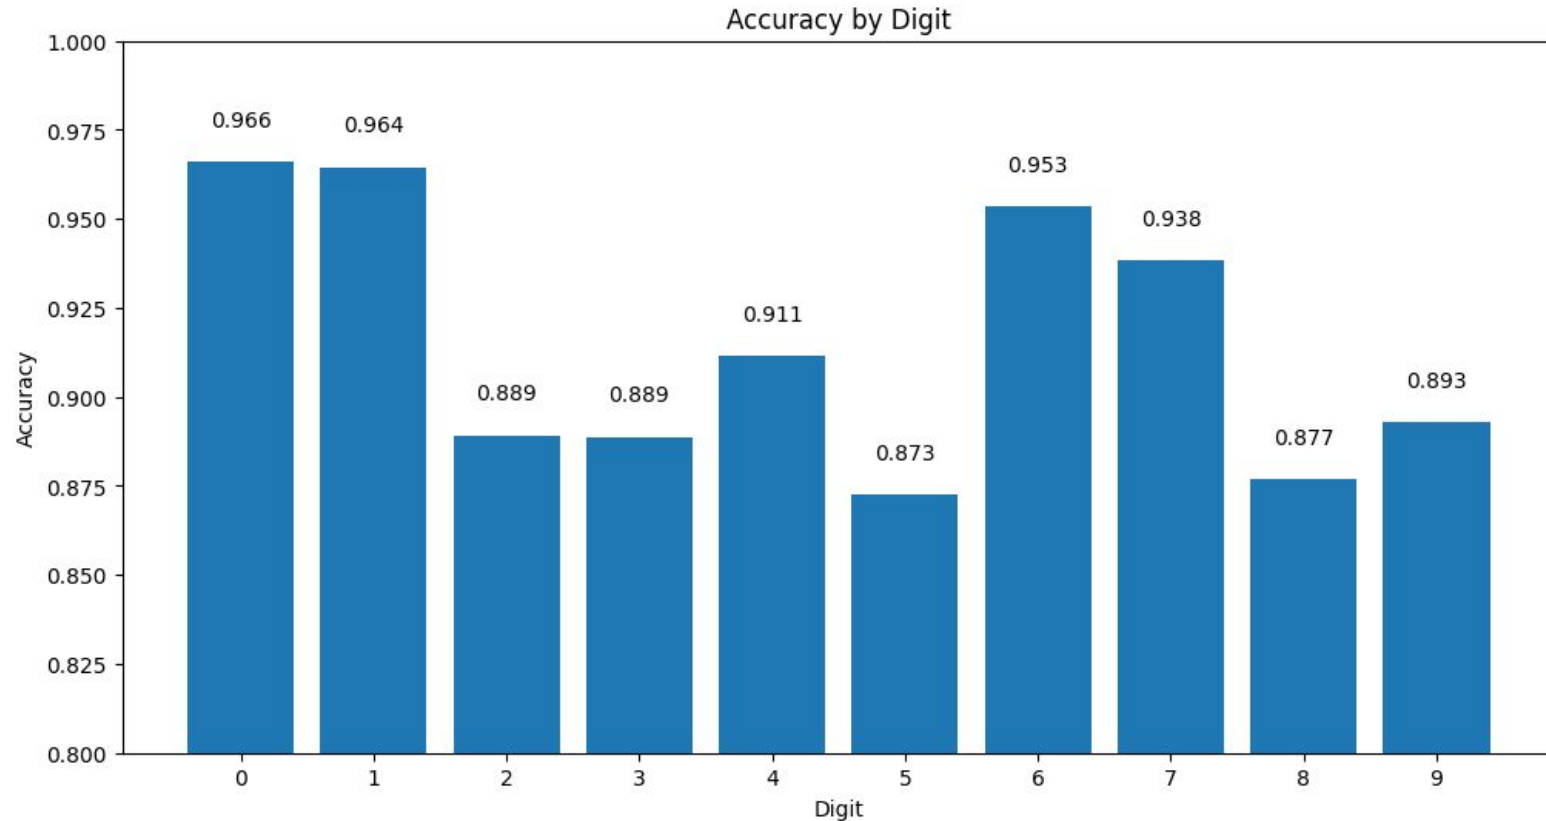
Sample predictions:



Logistic Regression - Confusion Matrix



Logistic Regression - Per Digit Performance



Logistic Regression - Strengths and Weaknesses

Strengths:

- good performance - hitting ~92% accuracy with minimal fuss
- fast training compared to deep learning approaches
- Super lightweight Easy to interpret what's happening under the hood
- Great baseline model for benchmarking more complex solutions

Weaknesses:

- Struggles with similar-looking digits (especially 3/8 and 4/9 pairs) for complex cases.
- Performance is lower than other complex approaches (CNNs easily hit 99%+)
- Might fail with real-world, messy handwriting unlike MNIST's clean dataset

SVM

- Tested with 3 kernels: linear, poly, sigmoid and 2 regularization values: 1, 10
- Evaluated on accuracy, precision, recall, and f1-score

SVM - Performance

Kernel: linear, C: 1
Mean CV Accuracy: 0.9050
Accuracy: 0.8990
Precision: 0.8993
Recall: 0.8990
F1-Score: 0.8985

Kernel: poly, C: 1
Mean CV Accuracy: 0.8870
Accuracy: 0.8925
Precision: 0.9067
Recall: 0.8925
F1-Score: 0.8952


Kernel: sigmoid, C: 1
Mean CV Accuracy: 0.9040
Accuracy: 0.8930
Precision: 0.8932
Recall: 0.8930
F1-Score: 0.8927

Kernel: linear, C: 10
Mean CV Accuracy: 0.9050
Accuracy: 0.8990
Precision: 0.8993
Recall: 0.8990
F1-Score: 0.8985

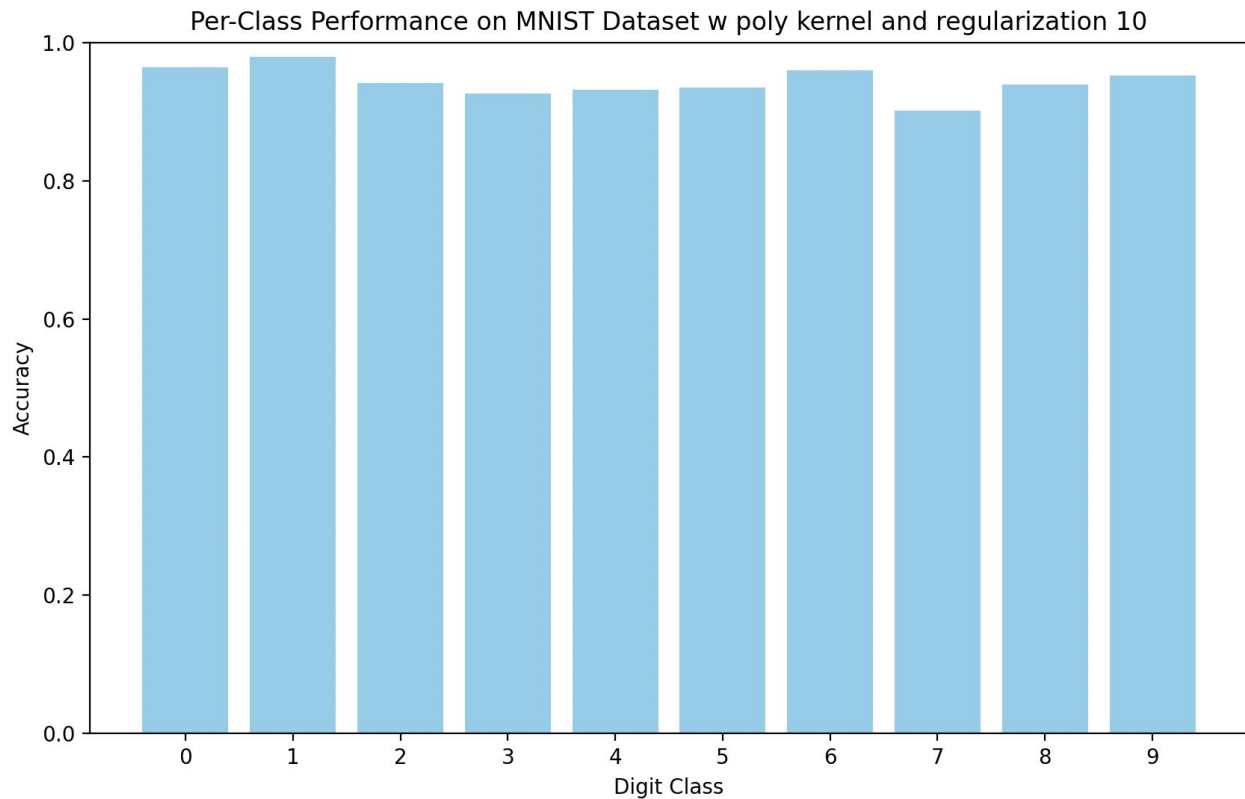
Kernel: poly, C: 10
Mean CV Accuracy: 0.9468
Accuracy: 0.9435
Precision: 0.9446
Recall: 0.9435
F1-Score: 0.9437

Kernel: sigmoid, C: 10
Mean CV Accuracy: 0.8686
Accuracy: 0.8615
Precision: 0.8622
Recall: 0.8615
F1-Score: 0.8611

SVM - Confusion Matrix



SVM - Per Digit Performance



SVM - Strengths and Weaknesses

Strengths:

- Good with high-dimensional data (like MNIST, 784)
- Different kernels for flexibility

Weaknesses:

- Slow training speed
- Complex parameter tuning needed
- Not naturally multi-class


KNN

- Perform 10-fold cross validation to find best k value from a preset range from 1-10
- KFold from sklearn to shuffle and split data into 10 folds
- For each data partition (one of the 10 folds), the model is evaluated using every k value
 - Accuracy scores obtained for each k are averaged across all partitions (folds)
- We repeat this process 10 times each time with a different seed which affects how the data is partitioned
 - Calculate the average of each k across the different seeds and find the k value with the highest score

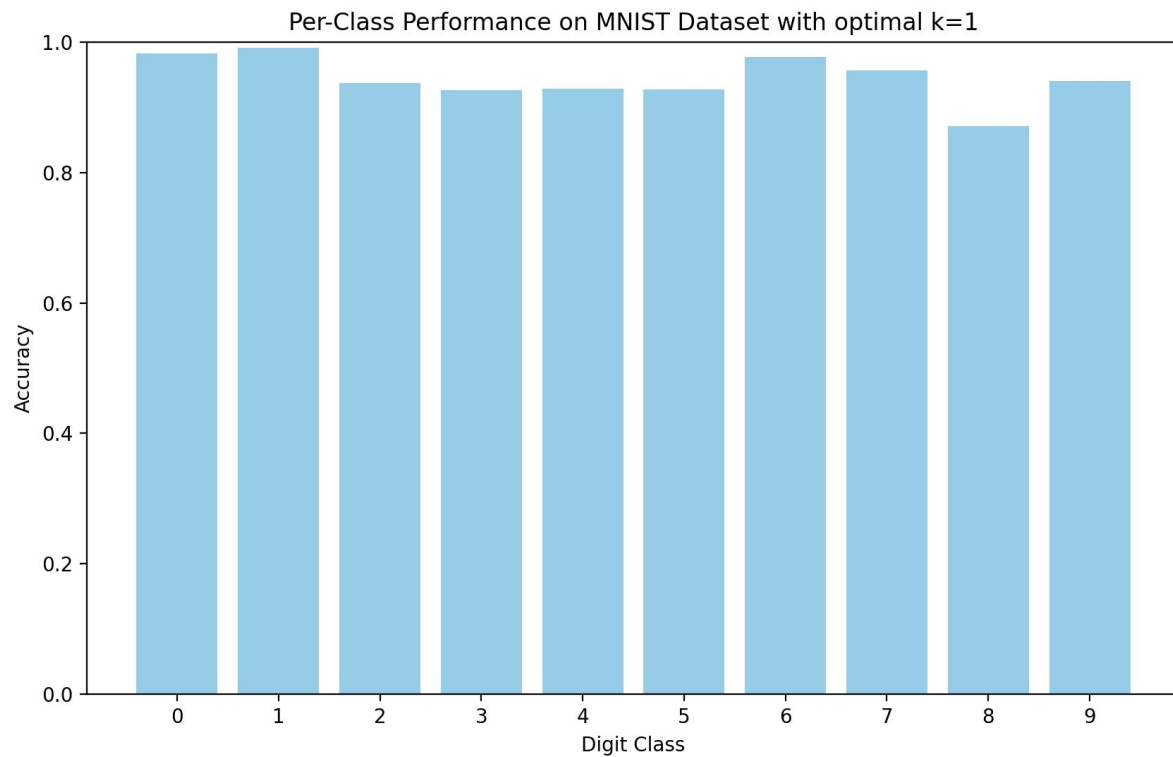
Performance

```
Overall average accuracy for k=1: 0.9446299999999999
Precision: 0.9452591178695622, Recall: 0.9435253235388925, F1 Score: 0.9436872645660201
Overall average accuracy for k=2: 0.9337799999999999
Precision: 0.9363279845984757, Recall: 0.9319876790556461, F1 Score: 0.9323037018713587
Overall average accuracy for k=3: 0.9441199999999999
Precision: 0.9456260438314821, Recall: 0.9428403119111625, F1 Score: 0.943240117027436
Overall average accuracy for k=4: 0.94331
Precision: 0.9453807930824262, Recall: 0.9420512976499419, F1 Score: 0.9426325733798849
Overall average accuracy for k=5: 0.9426200000000001
Precision: 0.9447239967984075, Recall: 0.9414358095104381, F1 Score: 0.9420058873588953
Overall average accuracy for k=6: 0.94141
Precision: 0.943811020361094, Recall: 0.9401930787505762, F1 Score: 0.9408527612020358
Overall average accuracy for k=7: 0.9394
Precision: 0.9418116629942419, Recall: 0.9382012952055359, F1 Score: 0.9387767395141526
Overall average accuracy for k=8: 0.9388000000000002
Precision: 0.9414181752975764, Recall: 0.9376079310699735, F1 Score: 0.9382526643522879
Overall average accuracy for k=9: 0.93814
Precision: 0.9409868902660732, Recall: 0.9369519339990571, F1 Score: 0.9376195658561464
Optimal k=1 metrics
Accuracy:0.9446299999999999
Precision: 0.9452591178695622
Recall: 0.9435253235388925
F1 Score: 0.9436872645660201
```

Confusion Matrix



Per Digit Performance



Strengths and Weaknesses

Strengths:

- Doesn't rely on obscure hyperparameters like weights and biases
 - Can get intuition on why a classification was made by looking at its nearest neighbors
- High accuracy on MNIST

Weaknesses:


- Computationally expensive to find the nearest neighbor
- On a high dimensional dataset like MNIST, could suffer from curse of dimensionality as distances get less meaningful with increase in more dimensions
 - However, this could be mitigated with dimensionality reduction such as PCA
- High memory usage as it has to store the entire training dataset

Kmeans Clustering


- Use elbow method to find number of clusters
- Initialize centroids using kmeans++
- Use KMeans from sklearn to cluster samples (784 dimensions)
- Use PCA to reduce samples and centroids to 2 dimensions and plot them
- Reconstruct flattened samples into original 28x28 shape and display the cluster means and some random samples in that cluster

Elbow Chart

Elbow Method for Optimal k




Elbow Method for Optimal k




Clusters



Clusters and Centroids with PCA



Clusters and Centroids with PCA



Visualized Results



Strengths and Weaknesses

Strengths:

- No labels required
- Low computation cost, especially with PCA

Weaknesses:

- Elbow method doesn't yield expected number of clusters
- Even in cases of MNIST where underlying structures of the data are known (number of classes), clustering is still inaccurate

Thank You

Questions?